



**RAJALAKSHMI ENGINEERING COLLEGE**

*Approved by AICTE | Affiliated to Anna University | Accredited by NAAC*

Department of Computer Science and Engineering

CS23334 Fundamentals of Data Science Lab


III semester II Year (2023R)

Name of the Student: **HEMALATHA ST**

Register Number: **240701189**

```
▶ import matplotlib.pyplot as plt  
years = [2018, 2019, 2020, 2021, 2022, 2023, 2024]  
india_scores = [285, 300, 295, 310, 325, 315, 330]  
pakistan_scores = [275, 290, 280, 295, 310, 305, 320]  
plt.figure(figsize=(4,2))  
plt.plot(years, india_scores, marker='o', color='blue', linewidth=2, label='India')  
plt.plot(years, pakistan_scores, marker='s', color='green', linewidth=2, linestyle='--', label='Pakistan')  
plt.title("Average Cricket Match Scores (India vs Pakistan)")  
plt.xlabel("Year")  
plt.ylabel("Average Team Score")  
plt.legend()  
plt.grid(True)  
plt.show()
```


Average Cricket Match Scores (India vs Pakistan)




```
▶ import matplotlib.pyplot as plt  
x=[1,2,3,4,5]  
y=[2,4,1,8,7]  
plt.figure(figsize=(4,4))  
plt.plot(x,y,color='blue',marker='o',linestyle='--',label='Line Plot')  
plt.title("Line Plot Example")  
plt.xlabel("X-axis")  
plt.ylabel("Y-axis")  
plt.legend()  
plt.grid(True)  
plt.show()
```

...


Line Plot Example




```
import matplotlib.pyplot as plt  
x=[45,23,67,35,67,88,67,89,56,54]  
y=[67,87,45,34,97,32,44,23,21,78]  
plt.figure(figsize=(5,4))  
plt.bar(x,y,color="green")  
plt.title("bar chart")  
plt.xlabel("fruits")  
plt.ylabel("quantity")  
plt.show()
```



```
import matplotlib.pyplot as plt
students = ['A', 'B', 'C', 'D', 'E', 'F']
math_marks = [78, 85, 69, 92, 74, 88]
science_marks = [82, 79, 75, 95, 70, 90]
plt.figure(figsize=(5,2))
plt.scatter(math_marks, science_marks, color='purple', marker='o', s=100)
plt.title("Students' Marks in Math vs Science")
plt.xlabel("Math Marks")
plt.ylabel("Science Marks")
plt.grid(True)
plt.show()
```



```
import pandas as pd
import matplotlib.pyplot as plt
df = pd.read_csv('airlines_flights_data.csv')
plt.hist(df['price'], edgecolor='black')
plt.xlabel('Price')
plt.ylabel('Number of Flights')
plt.title('Flight Price Distribution')
plt.show()
```



```
import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline

# Load penguins dataset
penguins = sns.load_dataset('penguins')

# Show first few rows
print(penguins.head())

# Set a seaborn style
sns.set_style("whitegrid")

# Distribution plot of flipper length with KDE - using 'coolwarm' palette
sns.displot(penguins['flipper_length_mm'].dropna(), kde=True, color='mediumvioletred')
plt.title('Distribution of Flipper Length (KDE)')
plt.show()

# Distribution plot without KDE - using color 'teal'
sns.displot(penguins['flipper_length_mm'].dropna(), kde=False, color='teal')
plt.title('Flipper Length Distribution (No KDE)')
plt.show()

# Jointplot between flipper length and body mass with custom colors
sns.jointplot(x='flipper_length_mm', y='body_mass_g', data=penguins, color='darkorange')
plt.show()
```

```
# Boxplot of body mass by species with palette
sns.boxplot(x='species', y='body_mass_g', data=penguins, palette=['#FF6347', '#4682B4', '#32CD32'])
plt.title('Body Mass Distribution by Species')
plt.show()

# Countplot of species with pastel colors
sns.countplot(x='species', data=penguins, palette='pastel')
plt.title('Count of Penguins by Species')
plt.show()


# Pie chart of species counts with custom colors
penguins['species'].value_counts().plot(
    kind='pie', autopct='%1.1f%%', startangle=140,
    colors=[ '#8A2BE2', '#00FA9A', '#1E90FF'],
    wedgeprops={'edgecolor': 'black'}
)
plt.title('Penguin Species Distribution')
plt.ylabel('')
plt.show()

# Bar chart of island counts with a dark palette
penguins['island'].value_counts().plot(kind='bar', color=[ '#D2691E', '#5F9EA0', '#9ACD32'])
plt.title('Number of Penguins by Island')
plt.xlabel('Island')
plt.ylabel('Count')
plt.show()
```

```
species      island bill_length_mm bill_depth_mm flipper_length_mm \
0  Adelie    Torgersen        39.1          18.7           181.0
1  Adelie    Torgersen        39.5          17.4           186.0
2  Adelie    Torgersen        40.3          18.0           195.0
3  Adelie    Torgersen        NaN            NaN           NaN
4  Adelie    Torgersen        36.7          19.3           193.0

body_mass_g      sex
0      3750.0   Male
1      3800.0 Female
2      3250.0 Female
3       NaN     NaN
4      3450.0 Female
```

Correlation Heatmap



```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

data = np.random.randint(1, 100, 20)
print("Original data:", data)

Q1 = np.percentile(data, 25)
Q3 = np.percentile(data, 75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

cleaned_data = data[(data >= lower_bound) & (data <= upper_bound)]

print("Cleaned data:", cleaned_data)

plt.figure(figsize=(12,5))

plt.subplot(1,2,1)
sns.histplot(data, bins=10, kde=True, color='skyblue')
plt.title('Histogram - Original Data')

plt.subplot(1,2,2)
sns.boxplot(x=data, color='lightgreen')
plt.title('Boxplot - Original Data')
```

```
plt.show()

plt.figure(figsize=(12,5))

plt.subplot(1,2,1)
sns.histplot(cleaned_data, bins=10, kde=True, color='orange')
plt.title('Histogram - Cleaned Data')


plt.subplot(1,2,2)
sns.boxplot(x=cleaned_data, color='lightcoral')
plt.title('Boxplot - Cleaned Data')

plt.show()
```


Original data: [74 44 91 30 29 78 29 22 54 95 15 67 48 15 98 91 33 25 45 20]

Cleaned data: [74 44 91 30 29 78 29 22 54 95 15 67 48 15 98 91 33 25 45 20]

Histogram - Original Data




Boxplot - Original Data



```
from sklearn.ensemble import RandomForestClassifier
from sklearn.datasets import load_iris
import pandas as pd
import matplotlib.pyplot as plt
iris = load_iris()
X = pd.DataFrame(iris.data, columns=iris.feature_names)
y = iris.target
model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X, y)
fi = pd.DataFrame({
    'Feature': X.columns,
    'Importance': model.feature_importances_
}).sort_values('Importance', ascending=False).head(10)
print(fi)
fi.plot.barh(x='Feature', y='Importance', legend=False, figsize=(6, 4), color='skyblue')
plt.gca().invert_yaxis()
plt.title("Top 10 Feature Importances")
plt.xlabel("Importance")
plt.ylabel("Feature")
plt.show()
```

	Feature	Importance
2	petal length (cm)	0.436130
3	petal width (cm)	0.436065
0	sepal length (cm)	0.106128
1	sepal width (cm)	0.021678

## Top 10 Feature Importances



```
import numpy as np
import pandas as pd
df = pd.read_csv('Salary_data.csv')
df.dropna(inplace=True)
print(df.info())
print(df.describe())
features = df.iloc[:, [0]].values
label = df.iloc[:, [1]].values
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(
    features, label, test_size=0.2, random_state=42
)
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(x_train, y_train)
print("Training Score:", model.score(x_train, y_train))
print("Testing Score:", model.score(x_test, y_test))
print("Coefficient:", model.coef_)
print("Intercept:", model.intercept_)
import pickle
pickle.dump(model, open('SalaryPred.model', 'wb'))
model = pickle.load(open('SalaryPred.model', 'rb'))
yr_of_exp = float(input("Enter Years of Experience: "))
yr_of_exp_NP = np.array([[yr_of_exp]])
Salary = model.predict(yr_of_exp_NP)
print(f"Estimated Salary for {yr_of_exp} years of experience is: {Salary[0][0]:.2f}")
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 29 entries, 0 to 28
```

```
Data columns (total 2 columns):
```

#	Column	Non-Null Count	Dtype
0	YearsExperience	29 non-null	float64
1	Salary	29 non-null	int64

```
dtypes: float64(1), int64(1)
```

```
memory usage: 596.0 bytes
```

```
None
```

	YearsExperience	Salary
count	29.000000	29.000000
mean	5.358621	76048.931034
std	2.877067	27983.440848
min	1.100000	37731.000000
25%	3.200000	56642.000000
50%	4.900000	63218.000000
75%	7.900000	101302.000000
max	10.500000	122391.000000

```
Training Score: 0.9393576307207374
```

```
Testing Score: 0.9170535247423002
```

```
Coefficient: [[9315.01199233]]
```

```
Intercept: [25125.45885762]
```

```
Enter Years of Experience: 4
```

```
Estimated Salary for 4.0 years of experience is: 62385.51
```

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
import numpy as np
iris = load_iris()
features = iris.data
label = iris.target
for i in range(1, 401):
    x_train, x_test, y_train, y_test = train_test_split(features, label, test_size=0.2, random_state=i)
    model = LogisticRegression(max_iter=200)
    model.fit(x_train, y_train)
    train_score = model.score(x_train, y_train)
    test_score = model.score(x_test, y_test)
    if test_score > train_score:
        print("Test: {:.3f} | Train: {:.3f} | Random state: {}".format(test_score, train_score, i))
x_train, x_test, y_train, y_test = train_test_split(features, label, test_size=0.2, random_state=42)
final_model = LogisticRegression(max_iter=200)
final_model.fit(x_train, y_train)
print("Train Accuracy:", final_model.score(x_train, y_train))
print("Test Accuracy:", final_model.score(x_test, y_test))
print(classification_report(label, final_model.predict(features), target_names=iris.target_names))
```

Train Accuracy: 0.975

Test Accuracy: 1.0

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	50
versicolor	1.00	0.94	0.97	50
virginica	0.94	1.00	0.97	50
accuracy			0.98	150
macro avg	0.98	0.98	0.98	150
weighted avg	0.98	0.98	0.98	150

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, classification_report
df = pd.read_csv('Iris1.csv')
df.info()
print(df['variety'].value_counts())
print(df.head())
features = df.iloc[:, :-1].values
label = df.iloc[:, 4].values
xtrain, xtest, ytrain, ytest = train_test_split(features, label, test_size=0.2, random_state=42)
model_KNN = KNeighborsClassifier(n_neighbors=5)
model_KNN.fit(xtrain, ytrain)
print("Training Accuracy:", model_KNN.score(xtrain, ytrain))
print("Testing Accuracy:", model_KNN.score(xtest, ytest))
print("\nConfusion Matrix:")
print(confusion_matrix(label, model_KNN.predict(features)))
print("\nClassification Report:")
print(classification_report(label, model_KNN.predict(features)))
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
 ---  -- 
 0   sepal_length  150 non-null   float64 
 1   sepal_width   150 non-null   float64 
 2   petal_length  150 non-null   float64 
 3   petal_width   150 non-null   float64 
 4   variety       150 non-null   object  
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
variety
Setosa      50
Versicolor  50
Virginica   50
Name: count, dtype: int64
    sepal_length  sepal_width  petal_length  petal_width  variety
0            5.1         3.5          1.4         0.2  Setosa
1            4.9         3.0          1.4         0.2  Setosa
2            4.7         3.2          1.3         0.2  Setosa
3            4.6         3.1          1.5         0.2  Setosa
4            5.0         3.6          1.4         0.2  Setosa
```

Training Accuracy: 0.9666666666666667

Testing Accuracy: 1.0

Confusion Matrix:

```
[[50  0  0]
 [ 0 47  3]
 [ 0  1 49]]
```

Classification Report:

	precision	recall	f1-score	support
Setosa	1.00	1.00	1.00	50
Versicolor	0.98	0.94	0.96	50
Virginica	0.94	0.98	0.96	50
accuracy			0.97	150
macro avg	0.97	0.97	0.97	150
weighted avg	0.97	0.97	0.97	150


```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans
data = {
    'CustomerID': range(1, 21),
    'Gender': ['Male', 'Female'] * 10,
    'Age': np.random.randint(18, 60, 20),
    'Annual Income (k$)': np.random.randint(15, 120, 20),
    'Spending Score (1-100)': np.random.randint(1, 100, 20)
}
df = pd.DataFrame(data)
print(df.info())
print(df.head())
features = df.iloc[:, [3, 4]].values
model = KMeans(n_clusters=5, random_state=42)
model.fit(features)
final_df = df.copy()
final_df['Label'] = model.predict(features)
sns.set_style("darkgrid")
palette = sns.color_palette("Spectral", 5)
plt.figure(figsize=(8, 6))
for label in range(5):
    cluster = final_df[final_df['Label'] == label]
    plt.scatter(cluster['Annual Income (k$)'], cluster['Spending Score (1-100)'],
                s=80, color=palette[label], label=f'Cluster {label+1}', alpha=0.7, edgecolor='black')
plt.title("Customer Segments (K-Means Clustering)", fontsize=14)
plt.xlabel("Annual Income (k$)")
plt.ylabel("Spending Score (1-100)")
```

```
plt.xlabel("Annual Income (k$)")
plt.ylabel("Spending Score (1-100)")
plt.legend()
plt.show()
wcss = []
for i in range(1, 10):
    model = KMeans(n_clusters=i, random_state=42)
    model.fit(features)
    wcss.append(model.inertia_)
plt.figure(figsize=(6, 4))
plt.plot(range(1, 10), wcss, marker='o')
plt.title("Elbow Method")
plt.xlabel("Number of Clusters")
plt.ylabel("WCSS")
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20 entries, 0 to 19
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   CustomerID      20 non-null     int64  
 1   Gender          20 non-null     object  
 2   Age             20 non-null     int64  
 3   Annual Income (k$) 20 non-null     int64  
 4   Spending Score (1-100) 20 non-null     int64  
dtypes: int64(4), object(1)
memory usage: 932.0+ bytes
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	24	91	22
1	2	Female	39	84	19
2	3	Male	58	101	59
3	4	Female	34	84	29
4	5	Male	37	32	26

Customer Segments (K-Means Clustering)



20

40

60


80

100

120

Annual Income (k\$)

## Elbow Method



```
from scipy.stats import chisquare

observed = [40, 25, 20, 15]
expected = [25, 25, 25, 25]

chi2_stat, p = chisquare(f_obs=observed, f_exp=expected)

print("Chi-Square Statistic:", chi2_stat)
print("P-Value:", p)

alpha = 0.05
if p < alpha:
    print("\nReject Null Hypothesis – observed distribution differs from expected.")
else:
    print("\nFail to Reject Null Hypothesis – no significant difference.")
```

Chi-Square Statistic: 14.0  
P-Value: 0.002905152774267437

Reject Null Hypothesis – observed distribution differs from expected.

```
import numpy as np
from scipy.stats import ttest_ind

group1 = np.array([12, 14, 15, 16, 14, 15, 13])
group2 = np.array([10, 11, 13, 12, 11, 10, 12])

t_stat, p = ttest_ind(group1, group2)
print("T-Statistic:", t_stat)
print("P-Value:", p)

alpha = 0.05
if p < alpha:
    print("\nReject Null Hypothesis - significant difference.")
else:
    print("\nFail to Reject Null Hypothesis - no significant difference.")
```

T-Statistic: 4.330127018922191

P-Value: 0.000978488712899117

Reject Null Hypothesis - significant difference.

```
import numpy as np
from statsmodels.stats.weightstats import ztest

group1 = np.array([50, 52, 49, 48, 47, 51, 50])
group2 = np.array([46, 47, 48, 45, 44, 46, 47])

z_stat, p = ztest(group1, group2)
print("Z-Statistic:", z_stat)
print("P-Value:", p)

alpha = 0.05
if p < alpha:
    print("\nReject Null Hypothesis - significant difference.")
else:
    print("\nFail to Reject Null Hypothesis - no significant difference.")
```

Z-Statistic: 4.1569219381653

P-Value: 3.225641456243845e-05

Reject Null Hypothesis - significant difference.

```
import numpy as np
from scipy.stats import f_oneway

group1 = np.array([23, 25, 27, 22, 24])
group2 = np.array([30, 31, 29, 32, 33])
group3 = np.array([40, 42, 41, 39, 43])

f_stat, p = f_oneway(group1, group2, group3)
print("F-Statistic:", f_stat)
print("P-Value:", p)

alpha = 0.05
if p < alpha:
    print("\nReject Null Hypothesis – at least one group differs.")
else:
    print("\nFail to Reject Null Hypothesis – all groups are similar.")
```

F-Statistic: 123.12643678160978

P-Value: 1.006506307348831e-08

Reject Null Hypothesis – at least one group differs.